

Sep 26, 18 21:59	mnist.py	Page 1/3
<pre>#!/bin/python3.6 #Ostap Voynarovskiy #CGML HW3 #Sept 26 2018 #Professor Curro  import numpy as np import tensorflow as tf import matplotlib.pyplot as plt import logging from tqdm import tqdm import os  #os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'  # Define globals and import dataset mnist = tf.keras.datasets.mnist (x_train,y_train),(x_test,y_test) = mnist.load_data() x_train,x_test = x_train/255.0 ,x_test/255.0      # go from 0-1 instead of 0-255 8 bit greyscale y_train = y_train.astype('int32')                # fix du e to tensorflow complaining y_test = y_test.astype('int32')                  # fix du e to tensorflow complaining  #import pdb; pdb.set_trace() # debug mode '''def cutData(f, cut):     l,w,h = f.shape     s,e = cut,w-cut #assumes square     fo = f[:,s:e,s]     print (fo.shape)     return fo ''' def cutData(f, cutx,cuty):     l,w,h = f.shape     sx,ex = cutx,w-cutx #assumes square     sy,ey = cuty,w-cuty #assumes square     fo = f[:,sy:ey,sx:ex]     print (fo.shape)     return fo  def genTrainAndVal(f,l): #split the features and labels of the training data 80: 20 train and validation     valPercent=20     lx,_,_=f.shape      z = f.shape[0]          # 60000 hopefully     s = np.arange(z)     np.random.shuffle(s)     fs = f[s]                # features shuffled     ls = l[s]                # labels shuffled     lx = f.shape[0]          # len of the features     nv = int( lx *.2)         # num validation samp     print (fs[nv:].shape, ls[nv:].shape, fs[:nv].shape, ls[:nv].shape)     return fs[nv:], ls[nv:], fs[:nv], ls[:nv]  #send this the output of only the first 2 returned val of genTrainAndVal def getBatch(feats,lab):     l,_,_ = feat.shape     choices = np.random.choice(l, size=BATCH_SIZE)     return feat[choices],lab[choices]  def cnn(features,labels,mode):     #layer 0      h,w = 20,20     L0 = tf.reshape(features["x"], [-1, h,w, 1]) #the -1 makes it guess how big its supposed to be</pre>		

Sep 26, 18 21:59	mnist.py	Page 2/3
<pre>is] # hyperParams k = 4 n = 5 l = 0.002 fil = 1 dr = .1 # conv layer 1 conv1 = tf.layers.conv2d(inputs=L0,filters=fil,kernel_size=[k, k], paddi ng="same", activation=tf.nn.elu) # pool layer 1 pool1 = tf.layers.max_pooling2d(inputs=conv1, pool_size=[n, n], strides= n) # flatten the layer for fully connected layer poolInt = int(w/n*h/n) * fil pFlat = tf.reshape(pool1, [-1, poolInt]) # fully connected layer #dense = tf.layers.dense(inputs=pFlat, units=32, activation=tf.nn.elu) # dropout layer dropout = tf.layers.dropout(inputs=pFlat, rate=dr, training=mode == tf.e stimator.ModeKeys.TRAIN) # logits layer layer that has the 10 outputs logits = tf.layers.dense(inputs=dropout, units=10)  predictions = {     "classes": tf.argmax(input=logits, axis=1),     "probabilities": tf.nn.softmax(logits, name="softmax_tensor")}  if mode == tf.estimator.ModeKeys.PREDICT:     return tf.estimator.EstimatorSpec(mode=mode, predictions=predicti ons)  num_params = np.sum([np.prod(v.get_shape().as_list()) for v in tf.traina ble_variables()]) print ("Num Params",num_params)  for v in tf.trainable_variables():     print (v.get_shape())  loss = tf.losses.sparse_softmax_cross_entropy(labels=labels, logits=logi ts) + 1*tf.reduce_sum([tf.nn.l2_loss(tV) for tV in tf.trainable_variables()])  if mode == tf.estimator.ModeKeys.TRAIN:     optimizer = tf.train.AdamOptimizer(learning_rate=0.001)     train_op = optimizer.minimize(loss=loss,global_step=tf.train.get _global_step())     return tf.estimator.EstimatorSpec(mode=mode, loss=loss, train_op =train_op)  eval_metric_ops = {     "accuracy": tf.metrics.accuracy(labels=labels, pre dictions=predictions["classes"])} return tf.estimator.EstimatorSpec(mode=mode, loss=loss, eval_metric_ops= eval_metric_ops)  def main():     mnist_classifier = tf.estimator.Estimator(model_fn=cnn)#,model_dir="./tm p/modelCheckpoint"     #tensors_to_log = {"probabilities": "softmax_tensor"}     #logging_hook = tf.train.LoggingTensorHook(tensors=tensors_to_log, every _n_iter=50)     tf.logging.set_verbosity(tf.logging.INFO)      train_input_fn = tf.estimator.inputs.numpy_input_fn(x={"x": tx},y=ty,bat ch_size=100,         num_epochs=None,shuffle=True)      mnist_classifier.train(input_fn=train_input_fn,steps=2000)#,hooks=[loggi</pre>		

Sep 26, 18 21:59

mnist.py

Page 3/3

```

ng_hook])

    eval_input_fn = tf.estimator.inputs.numpy_input_fn(x={"x": vx},y=vy,
        num_epochs=1,shuffle=False)

    eval_results = mnist_classifier.evaluate(input_fn=eval_input_fn)
    print(eval_results)

    #I did not have this here when I was tuning hyper Parameters
    eval_input_fn = tf.estimator.inputs.numpy_input_fn(x={"x": x_test},y=y_t
est,
        num_epochs=1,shuffle=False)

    eval_results = mnist_classifier.evaluate(input_fn=eval_input_fn)
    print(eval_results)

cutx = 4
cuty = 4
tx,ty,vx,vy = genTrainAndVal(x_train,y_train)
tx=cutData(tx,cutx,cuty)
vx=cutData(vx,cutx,cuty)
x_test = cutData(x_test,cutx,cuty)
print(tx.shape,'hi')
main()

#end of code

# just for shifts
def plotVal():
    len,_,_ = x_train.shape

    rn = np.random.randint(0,len-1)
    #rn = 26563
    print(len,rn)
    test_val=x_train[rn]

    fig1= plt.figure(1)
    dr=4 #pixels top bottom left and right to drop
    xc,yc = np.linspace(0+dr,27-dr,28-2*dr),np.linspace(27-dr,0+dr,28-2*dr)
    xv,yv = np.meshgrid(xc,yc)

    print(type(y_train))

    #reduce dimentions of the test data
    f,l = 0+dr, 28-dr
    z = test_val[f:l,f:l]

    CS = plt.contourf(xv,yv,z,cmap='gray')
    w= plt.xlabel('w')
    h= plt.ylabel('h')
    h.set_rotation(0)
    plt.title("MNIST")
    plt.axis('equal')
    #plt.clabel(CS, fontsize=9, inline=1)
    plt.show()

```

```
Num Params 187
(4, 4, 1, 1)
(1,)
(16, 10)
(10,)
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Starting evaluation at 2018-09-27-02:00:19
INFO:tensorflow:Graph was finalized.
2018-09-26 22:00:19.747814: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1484] Adding visible gpu devices: 0
2018-09-26 22:00:19.747856: I tensorflow/core/common_runtime/gpu/gpu_device.cc:965] Device interconnect StreamExecutor with strength 1 edge matrix:
2018-09-26 22:00:19.747861: I tensorflow/core/common_runtime/gpu/gpu_device.cc:971] 0
2018-09-26 22:00:19.747865: I tensorflow/core/common_runtime/gpu/gpu_device.cc:984] 0: N
2018-09-26 22:00:19.747952: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1097] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 2821 MB memory) -> physical GPU (device: 0, name: GeForce GTX 980, pci bus id: 0000:01:00.0, compute capability: 5.2)
INFO:tensorflow:Restoring parameters from /tmp/tmpf8hptawv/model.ckpt-2000
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Finished evaluation at 2018-09-27-02:00:19
INFO:tensorflow:Saving dict for global step 2000: accuracy = 0.8451, global_step = 2000, loss = 0.5841905
INFO:tensorflow:Saving 'checkpoint_path' summary for global step 2000: /tmp/tmpf8hptawv/model.ckpt-2000
{'loss': 0.5841905, 'accuracy': 0.8451, 'global_step': 2000}
(venv) ostap@ostap-All-Series:~/Documents/Deep Learning/hw3curro$
```