

Oct 04, 18 0:55

CIFAR10Res.py

Page 1/3

```
#!/bin/python3.5
# Ostap Voynarovskiy
# CGML HW4
# October 4 2018
# Professo Curro
import os
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import keras
from keras.models import Sequential
from keras.datasets import cifar10
from keras.preprocessing.image import ImageDataGenerator
from keras.callbacks import LearningRateScheduler
from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D, Activation, BatchNormalization, AveragePooling2D, Input
from keras import regularizers
from keras.models import Model
from keras.callbacks import ReduceLROnPlateau

from keras.optimizers import Adam
from keras.regularizers import l2

#from keras import backend as K

num_classes=10
BATCH_SIZE = 32
epochs = 200
DROP_RATE =.3
weight_decay = 1e-4
def genTrainAndVal(f,1): #split the features and labels of the training data 80:
20 train and validation
    lx=f.shape[0]
    z = f.shape[0]
    s = np.arange(z)
    np.random.shuffle(s)
    fs = f[s]
    ls = l[s]
    lx = f.shape[0]
    nv = int( lx *.2)
    print (fs[nv:].shape, ls[nv:].shape, fs[:nv].shape, ls[:nv].shape)
    return fs[nv:], ls[nv:], fs[:nv], ls[:nv]

datagen = ImageDataGenerator(featurewise_center=False,samplewise_center=False,
featurewise_std_normalization=False,samplewise_std_normalization=False,
zca_whitening=False,zca_epsilon=1e-06,rotation_range=0,width_shift_range
=0.1,
height_shift_range=0.1,shear_range=0.,zoom_range=0.,channel_shift_range=
0.,
fill_mode='nearest',cval=0.,horizontal_flip=True,vertical_flip=False,resca
le=None,
preprocessing_function=None,data_format=None,validation_split=0.0)

#modified from https://github.com/keras-team/keras/blob/master/examples/cifar10_
resnet.py
def lr_schedule(epoch):
    lr = 1e-3
    if epoch > 150:
        lr *= 0.5e-3
    elif epoch > 120:
        lr *= 1e-3
    elif epoch > 100:
        lr *= 1e-2
    elif epoch > 60:
        lr *= 1e-1
    print ('Learning rate: ', lr)
    return lr
```

Oct 04, 18 0:55

CIFAR10Res.py

Page 2/3

```
# load cifar 10
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
input_shape = x_train.shape[1:]
print(input_shape)

#convert to float and normalize
x_train,x_test = x_train.astype('float32'),x_test.astype('float32')
x_train,x_test = x_train/255,x_test/255

x_t,y_t,x_v,y_v =genTrainAndVal(x_train,y_train)

#print Shapes
print ("Training features shape: ", x_t.shape)
print ("Validation features shape: ",x_v.shape)
print ("Test features shape: ", x_test.shape)

#one hot encode the labels
y_t = keras.utils.to_categorical(y_t,num_classes)
y_v = keras.utils.to_categorical(y_v,num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
def res_layer(inputs, num_filters=16, kernel_size=3,
strides=1, activation='elu'):

    conv = Conv2D(num_filters,kernel_size=kernel_size,
strides=strides,padding='same',kernel_initializer='glorot_uniform',
kernel_regularizer=l2(weight_decay))

    x = inputs
    x1 = Conv2D(num_filters,kernel_size=kernel_size,
strides=strides,padding='same',kernel_initializer='glorot_uniform',
kernel_regularizer=l2(weight_decay))(x)
    x = BatchNormalization()(x1)
    x = Activation(activation)(x)
    x = Conv2D(num_filters,kernel_size=kernel_size,
strides=strides,padding='same',kernel_initializer='glorot_uniform',
kernel_regularizer=l2(weight_decay))(x) #32,32,16
    x = BatchNormalization()(x)
    x = Activation(activation)(x)
    x = Conv2D(num_filters,kernel_size=kernel_size,
strides=strides,padding='same',kernel_initializer='glorot_uniform',
kernel_regularizer=l2(weight_decay))(x)
    x = BatchNormalization()(x)
    x = keras.layers.add([x1,x])
    return x

def resNet(input_shape, num_classes=10):
    num_filters = 16
    print ("hi",input_shape)
    inputs = Input(shape=input_shape)
    print (inputs.shape)
    y = res_layer(inputs=inputs)

    for i in range(2):
        y = Activation("elu")(y)
        y = MaxPooling2D(pool_size=2, strides=2)(y)
        num_filters*=2
        y = res_layer(inputs=y, num_filters = num_filters )

    x = Activation("elu")(y)
    x = AveragePooling2D(pool_size=8)(x)
    y = Flatten()(x)
    outputs = Dense(num_classes,activation='softmax',kernel_initializer='he_nor
mal')(y)

    model = Model(inputs=inputs, outputs=outputs)
    return model
```

Oct 04, 18 0:55

CIFAR10Res.py

Page 3/3

```

model = resNet(input_shape, num_classes=10)

model.summary()
model.compile(loss="categorical_crossentropy",
              optimizer=keras.optimizers.Adam(lr=lr_schedule()),
              metrics=['accuracy'])

datagen.fit(x_t)

lr_scheduler = LearningRateScheduler(lr_schedule)

lr_reducer = ReduceLROnPlateau(factor=np.sqrt(0.1),
                               cooldown=0, patience=5, min_lr=0.5e-6)

callbacks = [lr_reducer, lr_scheduler]

model.fit_generator(datagen.flow(x_t, y_t, batch_size=BATCH_SIZE),
                   validation_data=(x_v, y_v),
                   steps_per_epoch=x_t.shape[0]/BATCH_SIZE,
                   epochs=epochs,
                   verbose=1,
                   callbacks=callbacks)

```

```

score = model.evaluate(x_test, y_test, verbose=1)

```

```

print("Test loss:", score[0])

```

```

print("Test accuracy:", score[1])

```

```

'''

```

I implemented a variation of the resnet20 from the paper <https://arxiv.org/pdf/1512.03385.pdf> and pulled some syntax for the callbacks and live data generation from the keras github implementation. https://github.com/keras-team/keras/blob/master/examples/cifar10_resnet.py I went through many different variations and found that some things just don't really make sense. For instance, when I ran this program with 30% dropout, I lost 2% accuracy on the test set. I realized however that the average pooling seemed to help. I also attempted a model that didn't use the skip connections. It's hard to say whether the skip connections helped or not, but they took a long time to figure out how to implement. The Cifar 100 model used the non skip connection version of the model and it seemed to work fine although it did have the benefit of top 5 accuracy. The learning rate drop however did add a lot of accuracy and was probably the most useful thing. It dropped multiple times and stabilized my model. While my final model was overfit, it achieved a better accuracy.

```

'''

```

Oct 04, 18 1:08

CIFAR10ResOUT.txt

Page 1/15

```
(venv) ostop@ostap-All-Series:~/Documents/DeepLearning/hw4curro$ python CIFAR10Res.py
Using TensorFlow backend.
(32, 32, 3)
(40000, 32, 32, 3) (40000, 1) (10000, 32, 32, 3) (10000, 1)
Training features shape: (40000, 32, 32, 3)
Validation features shape: (10000, 32, 32, 3)
Test features shape: (10000, 32, 32, 3)
hi (32, 32, 3)
(?, 32, 32, 3)
2018-10-04 00:18:06.742793: I tensorflow/core/platform/cpu_feature_guard.cc:141]
Your CPU supports instructions that this TensorFlow binary was not compiled to
use: AVX2 FMA
2018-10-04 00:18:06.819645: I tensorflow/stream_executor/cuda/cuda_gpu_executor.
cc:964] successful NUMA node read from SysFS had negative value (-1), but there
must be at least one NUMA node, so returning NUMA node zero
2018-10-04 00:18:06.820032: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1
411] Found device 0 with properties:
name: GeForce GTX 980 major: 5 minor: 2 memoryClockRate(GHz): 1.2785
pciBusID: 0000:01:00.0
totalMemory: 3.94GiB freeMemory: 3.05GiB
2018-10-04 00:18:06.870278: I tensorflow/stream_executor/cuda/cuda_gpu_executor.
cc:964] successful NUMA node read from SysFS had negative value (-1), but there
must be at least one NUMA node, so returning NUMA node zero
2018-10-04 00:18:06.870668: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1
411] Found device 1 with properties:
name: GeForce GTX 970 major: 5 minor: 2 memoryClockRate(GHz): 1.329
pciBusID: 0000:02:00.0
totalMemory: 3.94GiB freeMemory: 3.87GiB
2018-10-04 00:18:06.870818: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1
490] Adding visible gpu devices: 0, 1
2018-10-04 00:18:07.225328: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9
71] Device interconnect StreamExecutor with strength 1 edge matrix:
2018-10-04 00:18:07.225360: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9
77]      0 1
2018-10-04 00:18:07.225365: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9
90] 0:      N Y
2018-10-04 00:18:07.225368: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9
90] 1:      Y N
2018-10-04 00:18:07.225597: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1
103] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 wit
h 2758 MB memory) -> physical GPU (device: 0, name: GeForce GTX 980, pci bus id:
0000:01:00.0, compute capability: 5.2)
2018-10-04 00:18:07.247495: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1
103] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:1 wit
h 3599 MB memory) -> physical GPU (device: 1, name: GeForce GTX 970, pci bus id:
0000:02:00.0, compute capability: 5.2)
```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 32, 32, 3)	0	
conv2d_2 (Conv2D)	(None, 32, 32, 16)	448	input_1[0][0]
batch_normalization_1 (BatchNor	(None, 32, 32, 16)	64	conv2d_2[0][0]
activation_1 (Activation)	(None, 32, 32, 16)	0	batch_normaliza tion_1[0][0]

Oct 04, 18 1:08	CIFAR10ResOUT.txt		Page 2/15
conv2d_3 (Conv2D)	(None, 32, 32, 16)	2320	activation_1[0]
batch_normalization_2 (BatchNor	(None, 32, 32, 16)	64	conv2d_3[0][0]
activation_2 (Activation)	(None, 32, 32, 16)	0	batch_normaliza
conv2d_4 (Conv2D)	(None, 32, 32, 16)	2320	activation_2[0]
batch_normalization_3 (BatchNor	(None, 32, 32, 16)	64	conv2d_4[0][0]
add_1 (Add)	(None, 32, 32, 16)	0	conv2d_2[0][0]
activation_3 (Activation)	(None, 32, 32, 16)	0	batch_normaliza
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 16)	0	activation_3[0]
conv2d_6 (Conv2D)	(None, 16, 16, 32)	4640	max_pooling2d_1
batch_normalization_4 (BatchNor	(None, 16, 16, 32)	128	conv2d_6[0][0]
activation_4 (Activation)	(None, 16, 16, 32)	0	batch_normaliza
conv2d_7 (Conv2D)	(None, 16, 16, 32)	9248	activation_4[0]
batch_normalization_5 (BatchNor	(None, 16, 16, 32)	128	conv2d_7[0][0]
activation_5 (Activation)	(None, 16, 16, 32)	0	batch_normaliza
conv2d_8 (Conv2D)	(None, 16, 16, 32)	9248	activation_5[0]
batch_normalization_6 (BatchNor	(None, 16, 16, 32)	128	conv2d_8[0][0]
add_2 (Add)	(None, 16, 16, 32)	0	conv2d_6[0][0]
			batch_normaliza

Oct 04, 18 1:08	CIFAR10ResOUT.txt	Page 15/15
<pre> Learning rate: 5e-07 1250/1250 [=====] - 15s 12ms/step - loss: 0.2496 - acc: 0.9487 - val_loss: 0.4836 - val_acc: 0.8836 Epoch 191/200 Learning rate: 5e-07 1250/1250 [=====] - 15s 12ms/step - loss: 0.2496 - acc: 0.9500 - val_loss: 0.4843 - val_acc: 0.8836 Epoch 192/200 Learning rate: 5e-07 1250/1250 [=====] - 15s 12ms/step - loss: 0.2521 - acc: 0.9497 - val_loss: 0.4842 - val_acc: 0.8839 Epoch 193/200 Learning rate: 5e-07 1250/1250 [=====] - 15s 12ms/step - loss: 0.2526 - acc: 0.9476 - val_loss: 0.4828 - val_acc: 0.8841 Epoch 194/200 Learning rate: 5e-07 1250/1250 [=====] - 15s 12ms/step - loss: 0.2504 - acc: 0.9491 - val_loss: 0.4828 - val_acc: 0.8843 Epoch 195/200 Learning rate: 5e-07 1250/1250 [=====] - 14s 12ms/step - loss: 0.2496 - acc: 0.9490 - val_loss: 0.4835 - val_acc: 0.8841 Epoch 196/200 Learning rate: 5e-07 1250/1250 [=====] - 14s 12ms/step - loss: 0.2510 - acc: 0.9486 - val_loss: 0.4830 - val_acc: 0.8841 Epoch 197/200 Learning rate: 5e-07 1250/1250 [=====] - 15s 12ms/step - loss: 0.2524 - acc: 0.9494 - val_loss: 0.4837 - val_acc: 0.8842 Epoch 198/200 Learning rate: 5e-07 1250/1250 [=====] - 15s 12ms/step - loss: 0.2489 - acc: 0.9492 - val_loss: 0.4836 - val_acc: 0.8839 Epoch 199/200 Learning rate: 5e-07 1250/1250 [=====] - 15s 12ms/step - loss: 0.2506 - acc: 0.9485 - val_loss: 0.4855 - val_acc: 0.8829 Epoch 200/200 Learning rate: 5e-07 1250/1250 [=====] - 14s 12ms/step - loss: 0.2501 - acc: 0.9488 - val_loss: 0.4845 - val_acc: 0.8831 10000/10000 [=====] - 1s 105us/step Test loss: 0.5143441814422608 Test accuracy: 0.8784 </pre>		

Oct 03, 18 19:05	CIFAR100.py	Page 1/2
<pre>#!/bin/python3.5 # Ostop Voynarovskiy # CGML HW4 # October 4 2018 # Professo Curro import numpy as np import tensorflow as tf import matplotlib.pyplot as plt import keras from keras.models import Sequential from keras.datasets import cifar10 from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D, Activation, BatchNormalization from keras import regularizers from keras.metrics import top_k_categorical_accuracy #from keras import backend as K num_classes=100 BATCH_SIZE = 32 epochs = 32 #we achieve overfitting after like 15-20 epochs DROP_RATE =.5 weight_decay = 1e-4 def genTrainAndVal(f,l): #split the features and labels of the training data 80: 20 train and validation lx=f.shape[0] z = f.shape[0] s = np.arange(z) np.random.shuffle(s) fs = f[s] # features shuffled ls = l[s] # labels shuffled lx = f.shape[0] # len of the features nv = int(lx *.2) # num validation samp print (fs[nv:].shape, ls[nv:].shape, fs[:nv].shape, ls[:nv].shape) return fs[nv:], ls[nv:], fs[:nv], ls[:nv] # load cifar 10 #(x_train, y_train), (x_test, y_test) = cifar10.load_data() # load cifar 100 from keras.datasets import cifar100 (x_train, y_train), (x_test, y_test) = cifar100.load_data(label_mode='fine') #convert to float and normalize x_train,x_test = x_train.astype('float32'),x_test.astype('float32') x_train,x_test = x_train/255,x_test/255 x_t,y_t,x_v,y_v =genTrainAndVal(x_train,y_train) #print Shapes print ("Training features shape: ", x_t.shape) print ("Validation features shape: ",x_v.shape) print ("Test features shape: ", x_test.shape) #one hot encode the labels y_t = keras.utils.to_categorical(y_t,num_classes) y_v = keras.utils.to_categorical(y_v,num_classes) y_test = keras.utils.to_categorical(y_test, num_classes) model = Sequential() model.add(Conv2D(32,(4,4),padding='same',kernel_regularizer=regularizers.l2(weight_decay), data_format='channels_last', kernel_initializer='glorot_uniform', input_shape=x_t[0].shape)) model.add(Activation("elu")) model.add(BatchNormalization()) model.add(Conv2D(32,(3,3),strides=1,padding='same', kernel_regularizer=regulariz</pre>		

Oct 03, 18 19:05	CIFAR100.py	Page 2/2
<pre>ers.l2(weight_decay), kernel_initializer='glorot_uniform')) model.add(Activation("elu")) model.add(MaxPooling2D(pool_size=(2,2))) model.add(Dropout(DROP_RATE)) model.add(Conv2D(64,kernel_size=(3,3),padding='same',kernel_regularizer=regularizers.l2(weight_decay), data_format='channels_last', kernel_initializer='glorot_uniform')) model.add(Activation("elu")) model.add(BatchNormalization()) model.add(Conv2D(64,kernel_size=(4,4),padding='same',kernel_regularizer=regularizers.l2(weight_decay), data_format='channels_last', kernel_initializer='glorot_uniform')) model.add(Activation("elu")) model.add(BatchNormalization()) model.add(MaxPooling2D(pool_size=(2,2), strides=2)) model.add(Dropout(DROP_RATE)) model.add(Conv2D(128,kernel_size=(5,5),padding='same',kernel_regularizer=regularizers.l2(weight_decay), data_format='channels_last', kernel_initializer='glorot_uniform')) model.add(Activation("elu")) model.add(BatchNormalization()) model.add(Conv2D(128,kernel_size=(2,2),padding='same',kernel_regularizer=regularizers.l2(weight_decay), data_format='channels_last', kernel_initializer='glorot_uniform')) model.add(Activation("elu")) model.add(BatchNormalization()) model.add(MaxPooling2D(pool_size=(2,2), strides=2)) model.add(Dropout(.2)) model.add(Flatten()) model.add(Dense(num_classes,activation="softmax")) model.summary() model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=None, decay=0.0, amsgrad=False), metrics=["top_k_categorical_accuracy"]) model.fit(x_t, y_t, batch_size=BATCH_SIZE, epochs = epochs, verbose = 1, validation_data =(x_v,y_v)) score = model.evaluate(x_test,y_test,verbose= 1) print ("Test loss:", score[0]) print ("Test accuracy:", score[1])</pre>		

Oct 03, 18 19:13	CIFAR100OUT.txt	Page 1/4
<pre>(venv) ostop@ostap-All-Series:~/Documents/DeepLearning/hw4curro\$ python CIFAR100.py Using TensorFlow backend. (40000, 32, 32, 3) (40000, 1) (10000, 32, 32, 3) (10000, 1) Training features shape: (40000, 32, 32, 3) Validation features shape: (10000, 32, 32, 3) Test features shape: (10000, 32, 32, 3) 2018-10-03 18:52:53.455261: I tensorflow/core/platform/cpu_feature_guard.cc:141] Your CPU supports instructions that this TensorFlow binary was not compiled to use: AVX2 FMA 2018-10-03 18:52:53.548204: I tensorflow/stream_executor/cuda/cuda_gpu_executor. cc:964] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero 2018-10-03 18:52:53.548624: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1 411] Found device 0 with properties: name: GeForce GTX 980 major: 5 minor: 2 memoryClockRate(GHz): 1.2785 pciBusID: 0000:01:00.0 totalMemory: 3.94GiB freeMemory: 3.04GiB 2018-10-03 18:52:53.614475: I tensorflow/stream_executor/cuda/cuda_gpu_executor. cc:964] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero 2018-10-03 18:52:53.614870: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1 411] Found device 1 with properties: name: GeForce GTX 970 major: 5 minor: 2 memoryClockRate(GHz): 1.329 pciBusID: 0000:02:00.0 totalMemory: 3.94GiB freeMemory: 3.87GiB 2018-10-03 18:52:53.615025: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1 490] Adding visible gpu devices: 0, 1 2018-10-03 18:52:53.967614: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9 71] Device interconnect StreamExecutor with strength 1 edge matrix: 2018-10-03 18:52:53.967645: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9 77] 0 1 2018-10-03 18:52:53.967650: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9 90] 0: N Y 2018-10-03 18:52:53.967653: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9 90] 1: Y N 2018-10-03 18:52:53.967872: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1 103] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 wit h 2751 MB memory) -> physical GPU (device: 0, name: GeForce GTX 980, pci bus id: 0000:01:00.0, compute capability: 5.2) 2018-10-03 18:52:53.989668: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1 103] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:1 wit h 3599 MB memory) -> physical GPU (device: 1, name: GeForce GTX 970, pci bus id: 0000:02:00.0, compute capability: 5.2)</pre>		
Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 32)	1568
activation_1 (Activation)	(None, 32, 32, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_2 (Conv2D)	(None, 32, 32, 32)	9248
activation_2 (Activation)	(None, 32, 32, 32)	0
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 32)	0
dropout_1 (Dropout)	(None, 16, 16, 32)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	18496
activation_3 (Activation)	(None, 16, 16, 64)	0
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_4 (Conv2D)	(None, 16, 16, 64)	65600

Oct 03, 18 19:13	CIFAR100OUT.txt	Page 2/4
activation_4 (Activation)	(None, 16, 16, 64)	0
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_2 (Dropout)	(None, 8, 8, 64)	0
conv2d_5 (Conv2D)	(None, 8, 8, 128)	204928
activation_5 (Activation)	(None, 8, 8, 128)	0
batch_normalization_4 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_6 (Conv2D)	(None, 8, 8, 128)	65664
activation_6 (Activation)	(None, 8, 8, 128)	0
batch_normalization_5 (Batch Normalization)	(None, 8, 8, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_3 (Dropout)	(None, 4, 4, 128)	0
flatten_1 (Flatten)	(None, 2048)	0
dense_1 (Dense)	(None, 100)	204900
=====		
Total params: 572,068		
Trainable params: 571,236		
Non-trainable params: 832		
=====		
Train on 40000 samples, validate on 10000 samples		
Epoch 1/32		
40000/40000 [=====] - 16s 409us/step - loss: 4.1911 - t		
op_k_categorical_accuracy: 0.3479 - val_loss: 3.3574 - val_top_k_categorical_acc		
uracy: 0.4923		
Epoch 2/32		
40000/40000 [=====] - 15s 376us/step - loss: 3.1132 - t		
op_k_categorical_accuracy: 0.5538 - val_loss: 2.8131 - val_top_k_categorical_acc		
uracy: 0.6225		
Epoch 3/32		
40000/40000 [=====] - 15s 375us/step - loss: 2.6354 - t		
op_k_categorical_accuracy: 0.6573 - val_loss: 2.5696 - val_top_k_categorical_acc		
uracy: 0.6761		
Epoch 4/32		
40000/40000 [=====] - 15s 376us/step - loss: 2.3939 - t		
op_k_categorical_accuracy: 0.7151 - val_loss: 2.3879 - val_top_k_categorical_acc		
uracy: 0.7168		
Epoch 5/32		
40000/40000 [=====] - 15s 375us/step - loss: 2.2350 - t		
op_k_categorical_accuracy: 0.7537 - val_loss: 2.3999 - val_top_k_categorical_acc		
uracy: 0.7222		
Epoch 6/32		
40000/40000 [=====] - 15s 376us/step - loss: 2.1149 - t		
op_k_categorical_accuracy: 0.7805 - val_loss: 2.2855 - val_top_k_categorical_acc		
uracy: 0.7483		
Epoch 7/32		
40000/40000 [=====] - 15s 376us/step - loss: 2.0224 - t		
op_k_categorical_accuracy: 0.8023 - val_loss: 2.3183 - val_top_k_categorical_acc		
uracy: 0.7565		
Epoch 8/32		
40000/40000 [=====] - 15s 375us/step - loss: 1.9339 - t		
op_k_categorical_accuracy: 0.8208 - val_loss: 2.2308 - val_top_k_categorical_acc		
uracy: 0.7750		
Epoch 9/32		
40000/40000 [=====] - 15s 376us/step - loss: 1.8632 - t		
op_k_categorical_accuracy: 0.8364 - val_loss: 2.3028 - val_top_k_categorical_acc		
uracy: 0.7636		

Oct 03, 18 19:13	CIFAR100OUT.txt	Page 3/4
Epoch 10/32	40000/40000 [=====] - 15s 375us/step - loss: 1.8067 - t	
	op_k_categorical_accuracy: 0.8467 - val_loss: 2.2427 - val_top_k_categorical_acc	
	uracy: 0.7786	
Epoch 11/32	40000/40000 [=====] - 15s 376us/step - loss: 1.7510 - t	
	op_k_categorical_accuracy: 0.8609 - val_loss: 2.1793 - val_top_k_categorical_acc	
	uracy: 0.7912	
Epoch 12/32	40000/40000 [=====] - 15s 375us/step - loss: 1.7062 - t	
	op_k_categorical_accuracy: 0.8696 - val_loss: 2.2404 - val_top_k_categorical_acc	
	uracy: 0.7853	
Epoch 13/32	40000/40000 [=====] - 15s 376us/step - loss: 1.6602 - t	
	op_k_categorical_accuracy: 0.8784 - val_loss: 2.1732 - val_top_k_categorical_acc	
	uracy: 0.8046	
Epoch 14/32	40000/40000 [=====] - 15s 375us/step - loss: 1.6301 - t	
	op_k_categorical_accuracy: 0.8872 - val_loss: 2.2326 - val_top_k_categorical_acc	
	uracy: 0.7956	
Epoch 15/32	40000/40000 [=====] - 15s 376us/step - loss: 1.5979 - t	
	op_k_categorical_accuracy: 0.8934 - val_loss: 2.3233 - val_top_k_categorical_acc	
	uracy: 0.7844	
Epoch 16/32	40000/40000 [=====] - 15s 376us/step - loss: 1.5764 - t	
	op_k_categorical_accuracy: 0.8984 - val_loss: 2.2007 - val_top_k_categorical_acc	
	uracy: 0.8040	
Epoch 17/32	40000/40000 [=====] - 15s 375us/step - loss: 1.5455 - t	
	op_k_categorical_accuracy: 0.9033 - val_loss: 2.3343 - val_top_k_categorical_acc	
	uracy: 0.7897	
Epoch 18/32	40000/40000 [=====] - 15s 376us/step - loss: 1.5250 - t	
	op_k_categorical_accuracy: 0.9079 - val_loss: 2.2597 - val_top_k_categorical_acc	
	uracy: 0.8064	
Epoch 19/32	40000/40000 [=====] - 15s 376us/step - loss: 1.4950 - t	
	op_k_categorical_accuracy: 0.9141 - val_loss: 2.2843 - val_top_k_categorical_acc	
	uracy: 0.8049	
Epoch 20/32	40000/40000 [=====] - 15s 376us/step - loss: 1.4831 - t	
	op_k_categorical_accuracy: 0.9169 - val_loss: 2.3495 - val_top_k_categorical_acc	
	uracy: 0.7995	
Epoch 21/32	40000/40000 [=====] - 15s 376us/step - loss: 1.4702 - t	
	op_k_categorical_accuracy: 0.9209 - val_loss: 2.3060 - val_top_k_categorical_acc	
	uracy: 0.8093	
Epoch 22/32	40000/40000 [=====] - 15s 375us/step - loss: 1.4421 - t	
	op_k_categorical_accuracy: 0.9256 - val_loss: 2.3839 - val_top_k_categorical_acc	
	uracy: 0.7979	
Epoch 23/32	40000/40000 [=====] - 15s 376us/step - loss: 1.4391 - t	
	op_k_categorical_accuracy: 0.9255 - val_loss: 2.3491 - val_top_k_categorical_acc	
	uracy: 0.8038	
Epoch 24/32	40000/40000 [=====] - 15s 376us/step - loss: 1.4216 - t	
	op_k_categorical_accuracy: 0.9286 - val_loss: 2.3787 - val_top_k_categorical_acc	
	uracy: 0.7953	
Epoch 25/32	40000/40000 [=====] - 15s 376us/step - loss: 1.4010 - t	
	op_k_categorical_accuracy: 0.9313 - val_loss: 2.3989 - val_top_k_categorical_acc	
	uracy: 0.7959	
Epoch 26/32	40000/40000 [=====] - 15s 375us/step - loss: 1.3877 - t	
	op_k_categorical_accuracy: 0.9346 - val_loss: 2.3856 - val_top_k_categorical_acc	
	uracy: 0.8006	
Epoch 27/32		

Oct 03, 18 19:13	CIFAR100OUT.txt	Page 4/4
40000/40000 [=====] - 15s 377us/step - loss: 1.3773 - t		
op_k_categorical_accuracy: 0.9373 - val_loss: 2.3777 - val_top_k_categorical_acc		
uracy: 0.8066		
Epoch 28/32	40000/40000 [=====] - 15s 376us/step - loss: 1.3772 - t	
	op_k_categorical_accuracy: 0.9367 - val_loss: 2.4399 - val_top_k_categorical_acc	
	uracy: 0.8018	
Epoch 29/32	40000/40000 [=====] - 15s 377us/step - loss: 1.3589 - t	
	op_k_categorical_accuracy: 0.9395 - val_loss: 2.4524 - val_top_k_categorical_acc	
	uracy: 0.7953	
Epoch 30/32	40000/40000 [=====] - 15s 375us/step - loss: 1.3470 - t	
	op_k_categorical_accuracy: 0.9427 - val_loss: 2.4399 - val_top_k_categorical_acc	
	uracy: 0.8001	
Epoch 31/32	40000/40000 [=====] - 15s 377us/step - loss: 1.3399 - t	
	op_k_categorical_accuracy: 0.9427 - val_loss: 2.4711 - val_top_k_categorical_acc	
	uracy: 0.7989	
Epoch 32/32	40000/40000 [=====] - 15s 376us/step - loss: 1.3318 - t	
	op_k_categorical_accuracy: 0.9438 - val_loss: 2.5110 - val_top_k_categorical_acc	
	uracy: 0.7980	
10000/10000 [=====] - 1s 117us/step		
Test loss: 2.4556733959197996		
Test accuracy: 0.796		