

Oct 10, 18 21:11

ag.py

Page 1/2

```
#!/bin/python3.5
# Ostop Voynarovskiy
# CGML HW5
# October 9 2018
# Professo Curro

import keras
import numpy as np
import pandas as pd
from keras.models import Sequential
from keras.preprocessing.text import Tokenizer
from keras.layers import Embedding, Dense, Flatten, Conv1D, Dropout
from sklearn.model_selection import train_test_split
from keras.preprocessing.sequence import pad_sequences
from keras import regularizers
#import scikit
# Read Data from CSV files header = none to not ignore first row
train = pd.read_csv("./ag_news/train.csv", header=None)
test = pd.read_csv("./ag_news/test.csv", header=None)

#label the collumns
train.columns = ['cat', 'title', 'description']
test.columns = ['cat', 'title', 'description']

train['space'] = " "
train["text"] = train.title + train.space + train.description
test['space'] = " "
test["text"] = test.title + test.space + test.description

# split the data apart for val
train, val = train_test_split(train, test_size=.05)

# Create a Tokenizer instance
tokenizer = Tokenizer()
# Only fit on text this for training data
tokenizer.fit_on_texts(train.text)

# Create sequenced data for the vocabs
trainSeq = tokenizer.texts_to_sequences(train.text)
valSeq = tokenizer.texts_to_sequences(val.text)
testSeq = tokenizer.texts_to_sequences(test.text)

# Pad the vocabs so that they are all the same len
lenPad=185
trainPad= pad_sequences(trainSeq, maxlen=lenPad)
valPad= pad_sequences(valSeq, maxlen=lenPad)
testPad= pad_sequences(testSeq, maxlen=lenPad)

# how long to make the embedding vector
embeddingVectorLen = 32
vocabLength = len(tokenizer.word_index)

# one hot encode the labels
train_cat = np.array(train.cat-1)
train_cat = train_cat.reshape(train_cat.shape[0],1)
train_cat = keras.utils.to_categorical(train_cat, num_classes=4)

val_cat = np.array(val.cat-1)
val_cat = val_cat.reshape(val_cat.shape[0],1)
val_cat = keras.utils.to_categorical(val_cat, num_classes=4)

test_cat = np.array(test.cat-1)
test_cat = test_cat.reshape(test_cat.shape[0],1)
test_cat = keras.utils.to_categorical(test_cat, num_classes=4)

model = Sequential()
model.add(Embedding(vocabLength+1, embeddingVectorLen, input_length=lenPad))
#model.add(Conv1D(8, 32, padding = "same", kernel_regularizer=regularizers.l2()))
```

Oct 10, 18 21:11

ag.py

Page 2/2

```
#model.add(Conv1D(128, 32, padding = "same", dilation_rate=3, kernel_regularizer=regularizers.l2()))

model.add(Dropout(.7))
model.add(Flatten())
model.add(Dense(4, activation="softmax"))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adam(), metrics=['accuracy'])

model.fit(trainPad, train_cat,
          batch_size=64,
          epochs = 2,
          verbose = 1,
          validation_data = (valPad, val_cat))

score = model.evaluate(testPad, test_cat, verbose= 1)
print ("Test loss:", score[0])
print ("Test accuracy:", score[1])
```

Oct 10, 18 0:56

agOut.txt

Page 1/1

```

(venv) ostap@ostap-All-Series:~/Documents/DeepLearning/hw5curro$ CUDA_VISIBLE_DE
VICES=0 python ag.py Using TensorFlow backend.
Train on 114000 samples, validate on 6000 samples
Epoch 1/2
2018-10-10 00:51:42.447309: I tensorflow/core/platform/cpu_feature_guard.cc:141]
Your CPU supports instructions that this TensorFlow binary was not compiled to
use: AVX2 FMA
2018-10-10 00:51:42.510410: I tensorflow/stream_executor/cuda/cuda_gpu_executor.
cc:964] successful NUMA node read from SysFS had negative value (-1), but there
must be at least one NUMA node, so returning NUMA node zero
2018-10-10 00:51:42.510868: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1
411] Found device 0 with properties:
name: GeForce GTX 980 major: 5 minor: 2 memoryClockRate(GHz): 1.2785
pciBusID: 0000:01:00.0
totalMemory: 3.94GiB freeMemory: 2.77GiB
2018-10-10 00:51:42.510883: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1
490] Adding visible gpu devices: 0
2018-10-10 00:51:42.693621: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9
71] Device interconnect StreamExecutor with strength 1 edge matrix:
2018-10-10 00:51:42.693664: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9
77] 0
2018-10-10 00:51:42.693669: I tensorflow/core/common_runtime/gpu/gpu_device.cc:9
90] 0: N
2018-10-10 00:51:42.693824: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1
103] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 wit
h 2473 MB memory) -> physical GPU (device: 0, name: GeForce GTX 980, pci bus id:
0000:01:00.0, compute capability: 5.2)
114000/114000 [=====] - 9s 80us/step - loss: 0.4342 - a
cc: 0.8536 - val_loss: 0.2533 - val_acc: 0.9175
Epoch 2/2
114000/114000 [=====] - 8s 74us/step - loss: 0.2031 - a
cc: 0.9336 - val_loss: 0.2390 - val_acc: 0.9220
7600/7600 [=====] - 0s 27us/step
Test loss: 0.24406225435827908
Test accuracy: 0.9219736842105263
(venv) ostap@ostap-All-Series:~/Documents/DeepLearning/hw5curro$

```